

**From Summative Multiple-Choice Testing to Formative Practice:
Artificial Intelligence-Supported Question Generation,
Cognitive Analytics and Adaptive Pacing**

R. Bimberg, P. Kubica, K. Quibeldy-Cirkel

*Technische Hochschule Mittelhessen (GERMANY)
IU International University (GERMANY)*

Abstract: Multiple-choice (MC) exams are widely used in higher education but often provide little more than a summative score. This substantially limits their value for metacognitive reflection and self-regulated learning: students can see whether an answer is correct, yet they rarely understand which cognitive demands they reliably master and where targeted remediation should begin.

The paper presents *MC-Test*, an artificial intelligence (AI)-supported platform that strengthens MC testing by integrating item generation, feedback, and learner-facing analytics into a coherent architecture. A didactic data model structures each item with a learning objective, topic, weighting, and cognitive level, and deliberately constrains AI-generated questions to three operationalized levels (reproduction, application, and analysis). This constraint increases verifiability and reduces the risk of semantic errors without shifting to higher-order tasks (evaluation, creation) that are difficult to validate in an MC format.

To support formative use, items include a brief explanation, detailed rationales for all answer options, and a mini glossary for just-in-time term clarification; for more complex tasks, *MC-Test* adds step-by-step explanations. In a subsequent step, the system derives learning objectives per cognitive level from items and labels and provides them as a study resource (including glossary export). Based on these metadata, the platform produces learner-facing summaries: a topic view highlights domain-specific gaps and contextualizes metrics by showing how many items have been attempted, while a cognitive radar profile visualizes performance across the three levels and is displayed only when enough data are available to avoid overinterpretation.

As a second core contribution, *MC-Test* incorporates a layer of pedagogical control designed to reduce impulsive responding (e.g., rapid guessing). Dynamic lockout periods before submission support careful reading, and reflection windows after submission increase the likelihood that formative feedback is used. An adaptive “panic mode” disables pacing when the ratio of remaining time to remaining questions becomes critical, prioritizing fairness under time pressure. All pacing parameters are configurable at the course level and can be reduced or disabled to support accessibility accommodations.

Technically, the item pipeline uses a finite-state interaction model within the prompt architecture: configuration, internal planning, and strict JavaScript Object Notation (JSON) schema-conformant output improve consistency and parsability and facilitate downstream psychometric checks. For data protection and sovereignty, *MC-Test* runs in a containerized on-premise deployment and uses a local large language model (LLM) backend, so prompts and student data remain within institutional infrastructure; pseudonyms support anonymous participation. An initial usability study using the System Usability Scale (SUS) with a pilot cohort (N=20) yields a mean score of 70.38 (“OK”-Range per Bangor et al.), indicating good user acceptance. Next steps include randomized comparisons between conditions (or staged

rollout) and analyses of pacing-related behavior (rapid-guessing indicators), response times, label reliability, and item statistics. The application is open source and available under the MIT License on GitHub.

Keywords: Artificial Intelligence in Education, Multiple-Choice Question Generation, Student-Facing Learning Analytics, Adaptive Pacing, Privacy-Preserving Local Language Models, Formative Assessment.

1. INTRODUCTION

Multiple-choice (MC) questions are widely used in university assessment because they scale well, support objective scoring, and are easy to administer in large cohorts. However, typical MC platforms provide little more than a total score and item-level correctness, offering limited support for metacognitive reflection and strategic exam preparation. Learners often receive feedback that is too coarse to inform their study strategies: they know *that* they failed, but not *why*.

Recent advances in large language models (LLMs) enable not only the automated generation of MC items but also richer internal representations of their cognitive demands. When item metadata such as topic, difficulty, and cognitive level are made explicit, MC tests can be reinterpreted as data sources for learner-facing analytics. Simultaneously, LLM-based coding assistants can substantially accelerate programming tasks and lower the barrier for domain experts to prototype tailored educational tools [1]. At the same time, studies have started to assess the psychometric characteristics of LLM-generated MCQs, including difficulty and discrimination indices. While such items can exhibit competitive measurement properties in some settings, the findings also underline the necessity of systematic item-level validation rather than assuming quality from generation alone [13].

This paper positions *MC-Test* as a concrete instantiation of these trends. It is a web-based platform integrating:

- A didactic data model with explicit cognitive-level labels.
- An LLM-based pipeline using a finite state interaction model for generating valid items.
- Learner-facing analytics visualizing performance by cognitive level.
- Pedagogical control mechanisms (dynamic cooldowns and pacing) to mitigate rapid guessing.

The platform was initially implemented using commercial LLM APIs and was later migrated to a local LLM backend via Ollama. The work is guided by the following research questions:

- **RQ1:** How can cognitive-level classifications be operationalized in a student-facing assessment platform?
- **RQ2:** How can LLMs be embedded via robust prompt engineering to support psychometric quality checks and reliable item generation?
- **RQ3:** Which design considerations arise when migrating to a local LLM backend?

2. THEORETICAL BACKGROUND AND RELATED WORK

Cognitive Taxonomies and AI Validity: Bloom's taxonomy has long provided a framework for differentiating cognitive complexity. While LLMs excel at generating content for lower

taxonomic levels (remember, understand), recent studies indicate significant reliability issues when models attempt to generate tasks for higher-order evaluation or creation without human-in-the-loop verification [2]. Consequently, automated MC generation requires strict constraints to ensure semantic validity [3].

Student-Facing Learning Analytics & Self-Regulation: Student-facing learning analytics dashboards (SFLAD) aim to provide learners with interpretable visualizations to support self-regulated learning [4]. However, research suggests that pure data visualization is often insufficient; effective systems must provide active guidance or “nudges” to translate insights into behavioral change [5]. In addition, prior work provides design and evaluation recommendations for student-facing learning analytics dashboards [6]. *MC-Test* addresses this by combining analytics with active pacing mechanisms.

Privacy in Generative AI: The use of commercial LLM APIs in education raises significant privacy concerns regarding data leakage and GDPR compliance [7]. Recent work highlights the “privacy paradox” of LLMs, suggesting that on-premise deployment (local LLMs) is a strong mitigation option for sensitive educational contexts [8].

3. SYSTEM DESIGN: THE *MC-TEST* PLATFORM

3.1 Pedagogical Requirements

MC-Test was designed for formative assessment in STEM courses. Key requirements included:

- **Cognitive transparency:** Items must be classified by cognitive level.
- **Actionable feedback:** Explanations must link to learning objectives.
- **Exam readiness:** Support for time management skills.
- **Didactic control:** Educators must retain control over item quality.

3.2 System Architecture

The application is implemented using the Python framework Streamlit and has been migrated from a cloud-based prototype to a robust, containerized microservice architecture. The entire stack is orchestrated via Docker Compose and deployed on an on-premise institutional server to ensure full data sovereignty. Key components include:

- **Data Persistence:** A PostgreSQL database replaces lightweight file-based storage to handle concurrent user sessions and complex relational data (user profiles, audit logs) reliably.
- **Privacy-Compliant Rendering:** To eliminate data leakage to external Content Delivery Networks (CDNs), mathematical formulas are rendered using a self-hosted local **MathJax** instance.
- **Local Inference Engine:** A central architectural shift is the integration of Ollama as a local LLM backend. This allows the platform to execute open-weights models (e.g., DeepSeek) entirely within the university's infrastructure, ensuring that neither student data nor prompt logic leaves the secure environment.

3.3 Didactic Data Model

The core is a strict JSON data model, with a schema excerpt shown in Fig. 1. Key fields include:

- `topic` and `learning_objective`.
- `weight` (1–3) and `cognitive_level` (reproduction, application, analysis).
- `mini_glossary`: Contextual definitions to support just-in-time learning.
- `rationales`: Detailed explanations for why each option is correct or incorrect.

3.4 Cognitive-Level Taxonomy for AI-Generated MCQs

MC-Test adopts a restricted three-level taxonomy inspired by Bloom: 1. Reproduction (Weight 1): Recall of facts, definitions, or simple algorithms. 2. Application (Weight 2): Use of known concepts in slightly varied contexts. 3. Analysis (Weight 3): Interpretation of data, comparison, or inference. Crucially, levels 4–6 (Evaluation, Creation) are explicitly excluded. Our analysis aligns with findings that current models lack the semantic grounding to reliably validate complex evaluation tasks in a closed MC format [2] and with work emphasizing that automated MC generation requires strict constraints to ensure semantic validity [3]. Restricting the system to levels 1–3 increases the verifiability of generated items and rationales and can reduce the risk of semantic hallucinations; however, human review and empirical item analysis remain necessary to ensure factual and psychometric quality.

3.4.1 Operationalization and Label Reliability

Learner-facing analytics are predicated on cognitive-level labels; consequently, *MC-Test* operationalizes the three levels with a brief rubric and a verification step. In the present workflow, the following three steps are taken. First, instructors define the learning objective and the intended cognitive level. Second, the LLM is instructed to generate an item that matches the specified level and to provide rationales aligned with the learning objective. Third, items are reviewed prior to release. To assess the reliability of the labeling process, a double-coding procedure will be implemented. This procedure will be executed by two instructors who will utilize a predetermined rubric. The interrater agreement, as measured by metrics such as percent agreement and Cohen's kappa, will be reported as part of the evaluation. This approach serves to prevent over-precise interpretations of learner-facing visualizations that lack the foundation of stable labeling.

3.5 Prompt Architecture

Item generation is controlled by a context-engineered system instruction set that operates as an “Interactive MCQ Generator.” Instead of the traditional zero-shot prompting approach, a finite-state interaction model with a gated workflow is implemented: (1) Configuration Phase, where the model elicits all required parameters sequentially (e.g., topic, target audience, item count and weight distribution, option format, and optional curricular context) and requests explicit confirmation before generation; (2) Internal Blueprinting, where the model performs a hidden planning and consistency-check step (e.g., validating that difficulty weights map correctly to the difficulty profile and that option/answer constraints are satisfiable) without exposing intermediate reasoning in the output; and (3) Schema-Driven Output, where the response is constrained to a single, strictly parseable JSON object with explicit escaping and formatting rules (including stable `question_ids` and a `schema_version` for forward

compatibility), Fig. 1. To improve robustness across heterogeneous backends (cloud and local LLMs), the generated JSON is additionally subjected to automated schema and semantic validation, with an optional repair loop that prompts the model to apply minimal edits when violations are detected.

Fig. 1: JSON schema for question sets to be generated via AI web interface

```
{
  "meta": {
    "schema_version": "1.1",
    "title": "string (from Configuration Step 1)",
    "created": "DD.MM.YYYY HH:MM",
    "target_audience": "string (from Configuration Step 2)",
    "question_count": number,
    "difficulty_profile": {
      "easy": number,
      "medium": number,
      "hard": number
    },
    "time_per_weight_minutes": {
      "1": 0.5,
      "2": 0.75,
      "3": 1.0
    },
    "additional_buffer_minutes": 5,
    "test_duration_minutes": number
  },
  "questions": [
    {
      "question_id": "Q001",
      "question": "string (Must start with '1. ', '2. ' etc.)",
      "options": ["string", "string", "string"],
      "answer": number,
      "explanation": "string (2-4 sentences)",
      "weight": number,
      "topic": "string",
      "concept": "string",
      "cognitive_level": "string (Reproduction | Application | Analysis)",
      "extended_explanation": null OR {"title": "string", "steps": ["string"], "content": "string"},
      "mini_glossary": [
        {"term": "TermKey", "definition": "Definition string"}
      ]
    }
  ]
}
```

Post-production learning objectives: For each generated question set, the LLM also proposes a small set of topic labels (typically ~10) that structure the item bank. In a second, dedicated post-production prompt, *MC-Test* derives learning objectives per cognitive level from each item, its assigned topic, and its cognitive label. The resulting objectives—phrased with Bloom-aligned action verbs—are compiled into a learner-facing study resource that summarizes the intended learning targets across topics and levels.

4. LEARNER-FACING ANALYTICS AND PEDAGOGICAL CONTROL

4.1 Visualizing Topic Performance

To facilitate targeted remediation, the “Topic Performance” chart provides a macro-level visualization of learner competence across different subject domains, Fig. 2. Horizontal bars quantify the percentage of maximum points achieved per topic, allowing for the immediate identification of specific knowledge gaps. Crucially, the system contextualizes these absolute values by displaying the ratio of completed items to the total pool size in brackets (e.g., “3/5”). This annotation serves as a transparency mechanism: it signals to the learner that performance metrics derived from topics with a low volume of questions (small N) reduce

reliability and should be interpreted as a heuristic indicator rather than a conclusive assessment.




Fig. 2: Overview of learner competence across machine learning topics. Horizontal bars represent the percentage of maximum achievable points per subject domain, with absolute answer counts (correct/total) labeled below each bar.

4.2 Visualizing Cognitive Profiles

At the test level, MC-Test aggregates results into views by topic and cognitive level. A distinctive feature is the Cognitive Radar Chart, which visualizes the learner's profile across the three axes (Reproduction, Application, Analysis), Fig. 3. This exposes patterns such as “binge learning” (high reproduction, low analysis), nudging learners toward deeper understanding [9]. To facilitate metacognitive diagnosis, the Cognitive Radar Chart quantifies learner performance by calculating the ratio of achieved to potential points across the three taxonomic levels. This visualization allows for the identification of distinct cognitive archetypes based on the morphological characteristics of the resulting plot. To reduce over-interpretation, MC-Test treats the radar profile as a heuristic indicator rather than a diagnostic measurement. The chart is only displayed when a minimum number of items per axis is available; otherwise, the interface flags “insufficient data” and recommends additional practice before interpreting the profile. For example, a ‘Theorist’ or ‘Crammer’ profile—characterized by a singular spike in Reproduction with steep drop-offs in higher-order levels—indicates superficial rote learning with limited transfer capabilities, prompting the system to recommend application-oriented tasks. Conversely, a ‘Practitioner’ profile (high Application, low Reproduction) suggests procedural fluency masked by theoretical deficits, triggering recommendations to revisit foundational concepts. Other patterns include the ‘Analyst’ (strong problem-solving with foundational gaps) and the ‘Zigzag’ (inconsistent performance), enabling the platform to derive specific, actionable interventions for each learner type, provided sufficient item volume ensures statistical reliability.




Fig. 3: Cognitive radar charts summarizing learner performance across the three cognitive levels (reproduction, application, and analysis) for three learner archetypes: Crammer (left), Practitioner (centre), and Theorist (right) (shown only when sufficient items per axis are available).

4.2 Feedback and Explanations

For complex items, explanations are extended with step-by-step reasoning. To support vocabulary acquisition, a two-tiered glossary concept is used: a mini-glossary field for immediate help and a summative PDF glossary for post-test review.

4.3 Pedagogical Control: Cooldowns and Panic Mode

A unique feature of *MC-Test* is its active management of learner behavior to encourage reading and reflection and to deter impulsive responding (e.g., rapid guessing). Rather than claiming to “enforce” deep processing, the platform implements configurable pacing nudges whose behavioral effects are treated as empirical questions to be evaluated. The system implements a Pedagogical Control layer consisting of three mechanisms: Pre-Answer Cooldown (reading support): to increase the likelihood that students read the question before responding, the “Submit” button remains disabled for a dynamic duration upon loading a new item. This duration is calculated based on the item’s weight (complexity) and text length. Post-Answer Cooldown (Reflection Support): After submitting an answer, navigation to the next question is delayed. This creates a mandatory window for reviewing the formative feedback [10]. The duration scales with the complexity of the explanation (e.g., +10s for standard, +20s for extended explanations). Adaptive Fairness (“Panic Mode”): While pacing can support metacognitive reflection, it can be detrimental under time pressure. The system continuously monitors the ratio of remaining time to remaining questions. If this drops below a critical threshold (e.g., <15 seconds per question), the system activates “Panic Mode”. In this state, all cooldowns are overridden to allow the student to finish the attempt, prioritizing fairness over didactic constraint in critical moments. Accessibility and accommodations: Because fixed pacing can disadvantage some learners (e.g., assistive-technology users or students with documented accommodations), all cooldown parameters are configurable at the course level and can be disabled or adjusted where required. Panic Mode further acts as a fairness safeguard under time pressure, prioritizing completion over pacing constraints.

5. IMPLEMENTATION AND MIGRATION TO LOCAL LLM BACKEND

The platform was migrated from commercial cloud APIs to an institutional server with a local LLM backend (Ollama) to address data protection and sustainability.

- **Prompt Refactoring:** The system prompt was adapted to local model characteristics, prioritizing JSON validity and robustness.
- **Privacy:** Local deployment ensures that no student data leaves the institution, adhering to “Privacy by Design” principles [7], [8].
- **Sustainability:** On-premise hosting mitigates the costs and rate limits associated with commercial API usage.
- **Data minimization and logging:** The deployment is designed to minimize sensitive data exposure. Only the information required for assessment and analytics is stored (e.g., answers, timestamps, aggregate scores), plus minimal account metadata and security audit logs needed to operate the system reliably. Prompts are constructed to avoid personal identifiers, and the local setup ensures that student data and prompt logic do not leave the secure environment. Access to logs and databases is restricted to authorized staff, and retention policies can be applied to limit long-term storage.
- **Anonymity and pseudonyms:** *MC-Test* is designed for anonymous participation. Users do not create personal accounts, and no real names are collected. At the start of a session, learners select a pseudonym from a predefined list of Nobel and Turing Award laureate names; this identifier is used only to keep session data consistent (e.g., to display progress and aggregate analytics) without enabling real-world identification of participants, Fig. 5.
- **Open-source availability:** To support transparency and reproducibility, the *MC-Test* Streamlit application is released under the MIT License and publicly available on GitHub (<https://github.com/kqc-real/streamlit>) The repository includes documentation for deployment and configuration (e.g., Docker Compose) as well as the schema-driven artifacts used in the local inference pipeline.

6. EVALUATION DESIGN AND OUTLOOK

6.1 Planned Study Design

The evaluation consists of pilot deployments in STEM modules. Instruments include system log data (pacing choices, cognitive profiles, response times), pre- and post-questionnaires on metacognitive awareness and perceived usefulness, and interviews. A key next step is a comparative study design (e.g., A/B testing or within-course phased rollout) to test whether the Pedagogical Control features measurably reduce rapid-guessing indicators and whether they affect learning outcomes or study strategies. The planned evaluation also includes checks of the stability of cognitive-level labeling (double-coding) and basic item-level statistics to support psychometric interpretation of generated question sets.

Initial informal trials indicate appreciation for the cognitive-level breakdown and the transparency of AI-generated rationales. Reactions to the pacing helper were mixed, highlighting the need for the adaptive “Panic Mode” which was subsequently implemented. These observations are exploratory and will be complemented by controlled comparisons using behavioral log metrics and learner-reported measures.

6.2 Preliminary Results: System Usability

To validate the user experience (UX) and technical acceptance of the platform, a standardized usability evaluation was conducted using the System Usability Scale (SUS) [11]. The survey was administered to a pilot cohort of 20 participants. The analysis yielded a mean SUS score

of 70.38. According to the adjective rating scale defined by Bangor et al. [12], this score places the *MC-Test* platform in the “OK” category, indicating good user acceptance despite the complexity of the integrated pedagogical features. Fig. 4 illustrates the distribution of individual participant scores. As shown in the diagram, the responses are plotted against the interpretive acceptability ranges (e.g., “Acceptable,” “Good,” “Excellent”). The visualization reveals a left-skewed distribution, confirming that the majority of users perceived the system as easy to use. Importantly, SUS primarily captures perceived usability; it does not establish learning benefits or behavioral effects. These aspects will be examined in subsequent studies using log-based rapid-guessing indicators, learner feedback on fairness/autonomy, and (where feasible) learning outcome measures.




Fig. 4: Distribution of System Usability Scale (SUS) scores ($N=20$); background bands reflect the adjective rating scale [12], the distribution is left-skewed (skewness = -0.625), indicating that higher usability scores are more common with some lower outliers.

6.3 Limitations

This paper primarily contributes a system design and preliminary usability evidence. The proposed benefits of cognitive profiling and pacing support (e.g., reduced rapid guessing, improved metacognitive regulation) require empirical validation beyond SUS. Cognitive-level labels can be interpretive; therefore, reliability checks (double-coding) and psychometric item analyses are necessary to avoid overconfident interpretations of learner-facing analytics. Finally, pacing mechanisms may affect learner autonomy and accessibility; configuration and accommodation pathways are required to avoid unintended disadvantages.

Recent studies explicitly assess the psychometric properties of LLM-generated questions, reinforcing the need to report item-level statistics when using AI-generated MCQs in practice [13]. Beyond classical item analysis, future work may also leverage “generative student” simulations—LLM-based profiles that approximate response patterns across different knowledge states—to flag ambiguous or low-discrimination items before deployment [14].

7. Conclusion

MC-Test demonstrates that effective AI-supported assessment requires more than simply calling an API. By combining engineered system instructions (Finite-State Interaction Model and schema-constrained generation) with didactic data models and pedagogical pacing support (Cooldowns/Panic Mode), the platform aims to transform the MCQ format from a summative drill into a formative, metacognitive tool. The restriction to verifiable cognitive levels (Bloom 1–3) and the migration to local LLMs strengthen the platform’s verifiability and data-sovereignty posture while supporting robust, parseable outputs.

The primary contributions of this work are: (1) a didactic item data model linking learning objectives, cognitive levels, and formative rationales; (2) an operational concept of Pedagogical Control (configurable cooldowns plus Panic Mode) intended to nudge reading and reflection under realistic time constraints; and (3) an architecture for local, schema-constrained LLM inference in an assessment context. Initial usability results are encouraging; however, evidence for learning impact, behavioral effects (e.g., rapid-guessing reduction), label reliability, and psychometric item quality will be addressed in planned comparative and log-based evaluations.




Fig. 5: *MC-Test*'s interface showing end-of-session feedback and the Topic Performance dashboard.

REFERENCES

- [1] T. Weber, M. Brandmaier, A. Schmidt, and S. Mayer, "Significant Productivity Gains through Programming with Large Language Models," *Proc. ACM Hum.-Comput. Interact.*, vol. 8, EICS, Art. 256, 2024, doi:10.1145/3661145.
- [2] N. Scaria, S. D. Chenna, and D. Subramani, "Automated Educational Question Generation at Different Bloom's Skill Levels Using Large Language Models: Strategies and Evaluation," in *Artificial Intelligence in Education (AIED 2024)*, LNAI vol. 14830, pp. 165–179, 2024, doi: 10.1007/978-3-031-64299-9_12.
- [3] P. Stalder, et al., "Ensuring Quality in AI-Generated Multiple-Choice Questions for Higher Education with the QUEST Framework," in *Communications in Computer and Information Science*, Springer, 2024.
- [4] R. Bodily and K. Verbert, "Trends and issues in student-facing learning analytics reporting systems research," in *Proc. 7th Int. Conf. Learning Analytics & Knowledge (LAK '17)*, Vancouver, BC, Canada, Mar. 2017, pp. 309–318, doi: 10.1145/3027385.3027403.
- [5] L. de Vreugd, et al., "Learning Analytics Dashboard Design and Evaluation to Support Student Self-Regulation of Study Behavior," *Journal of Learning Analytics*, vol. 11, no. 3, 2024.
- [6] I. Jivet, M. Scheffel, M. Specht, and H. Drachsler, "License to evaluate: Preparing learning analytics dashboards for educational practice," in *Proc. 8th Int. Conf. Learning Analytics & Knowledge (LAK '18)*, Sydney, NSW, Australia, Mar. 2018, pp. 31–40, doi: 10.1145/3170358.3170421.
- [7] H. Drachsler and W. Grelle, "Privacy and Analytics: It's a DELICATE Issue—A Checklist for Trusted Learning Analytics," in *Proceedings of the Sixth International Conference on Learning Analytics & Knowledge*, Edinburgh, UK, 2016, pp. 89–98.
- [8] Y. Shanmugarasa, M. Ding, M. A. P. Chamikara, and T. Rakotoarivelo, "SoK: The Privacy Paradox of Large Language Models: Advancements, Privacy Risks, and Mitigation," arXiv preprint arXiv:2506.12699, 2025.
- [9] Z. Yan and D. Carless, "Self-Assessment Is About More Than Self: The Enabling Role of Feedback Literacy," *Assessment & Evaluation in Higher Education*, vol. 47, no. 7, pp. 1116–1128, 2022.
- [10] N. E. Winstone and D. Carless, *Designing Effective Feedback Processes in Higher Education: A Learning-Focused Approach*. London: Routledge, 2019.
- [11] J. Brooke, "SUS: A 'Quick and Dirty' Usability Scale," in *Usability Evaluation in Industry*, P. W. Jordan, B. Thomas, B. A. Weerdmeester, and I. L. McClelland, Eds. London, U.K.: Taylor & Francis, 1996, pp. 189–194.
- [12] A. Bangor, P. T. Kortum, and J. T. Miller, "Determining what individual SUS scores mean: Adding an adjective rating scale," *Journal of Usability Studies*, vol. 4, no. 3, pp. 114–123, 2009.
- [13] S. Bhandari, Y. Liu, Y. Kwak, and Z. A. Pardos, "Evaluating the psychometric properties of ChatGPT-generated questions," *Computers and Education: Artificial Intelligence*, vol. 7, Art. 100284, 2024, doi: 10.1016/j.caeari.2024.100284.
- [14] X. Lu and X. Wang, "Generative Students: Using LLM-Simulated Student Profiles to Support Question Item Evaluation," in *Proc. 11th ACM Conf. on Learning @ Scale (L@S '24)*, 2024, pp. 16–27, doi: 10.1145/3657604.3662031.
- [15]